




AI Conclave

Online






AIOT201

Improve resiliency using Amazon Aurora DSQL, multi- region serverless distributed SQL database

Nethravathi Muddarajaiah

Principal Database Solutions Architect
AWS India



Agenda

- **What is Amazon Aurora DSQL?**
- **Architecture**
- **Cross-region and scalability**
- **Demo**



“The first principle of successful scalability is to batter the consistency mechanisms down to a minimum, move them off the critical path, hide them in a rarely visited corner of the system, and then make it as hard as possible for application developers to get permission to use them”

James Hamilton

Amazon SVP/Distinguished Engineer



Traditional relational databases

CHALLENGES

Scale

Capacity limits of traditional databases results in significant effort to right-size workloads

Availability

Lower resiliency leads to unplanned downtimes impacting database availability

Functional challenges

Management overhead

Infrastructure management, patching, upgrades requires high engineering effort for maintenance

Complexity

Infrastructure setup, database installation, configuration, tuning and management is complex and high effort

Operational challenges



Amazon Aurora DSQL is...



Amazon Aurora DSQL - Fastest distributed SQL database

AVAILABLE IN PUBLIC PREVIEW

- A relational SQL database optimized for transactional workloads
- Virtually unlimited scalability
- No infrastructure management
- 99.999% multi-region availability
- Automatic failure recovery
- PostgreSQL compatible





Amazon Aurora DSQL

CLOUD-NATIVE, SERVERLESS DISTRIBUTED SQL DATABASE WITH VIRTUALLY UNLIMITED SCALABILITY AND HIGHEST AVAILABILITY



Virtually unlimited scaling

Instantly scales compute, read, write and storage resources independently, both up and down



Business continuity

99.999% availability with active-active access in a multi-region cluster



Serverless

No servers to provision, patch, or manage, and no software to install, maintain, or operate.



Fast and Easy

Create and query a database in less than a minute with familiarity of PostgreSQL





Amazon Aurora DSQL

CLOUD-NATIVE, SERVERLESS DISTRIBUTED SQL DATABASE WITH VIRTUALLY UNLIMITED SCALABILITY AND HIGHEST AVAILABILITY



Virtually unlimited scaling

Instantly scales compute, read, write
and storage resources independently,
both up and down





Amazon Aurora DSQL

CLOUD-NATIVE, SERVERLESS DISTRIBUTED SQL DATABASE WITH VIRTUALLY UNLIMITED SCALABILITY AND HIGHEST AVAILABILITY



Business Continuity

99.999% availability with active-active access in a multi-region cluster





Amazon Aurora DSQL

CLOUD-NATIVE, SERVERLESS DISTRIBUTED SQL DATABASE WITH VIRTUALLY UNLIMITED SCALABILITY AND HIGHEST AVAILABILITY



Serverless

No servers to provision, patch, or manage, and no software to install, maintain, or operate.





Amazon Aurora DSQL

CLOUD-NATIVE, SERVERLESS DISTRIBUTED SQL DATABASE WITH VIRTUALLY UNLIMITED SCALABILITY AND HIGHEST AVAILABILITY



Fast and Easy

Create and query a database in
less than a minute with
familiarity of PostgreSQL



© 2025, Amazon Web Services, Inc. or its affiliates. All rights reserved. Amazon Confidential and Trademark.



Amazon Aurora DSQL

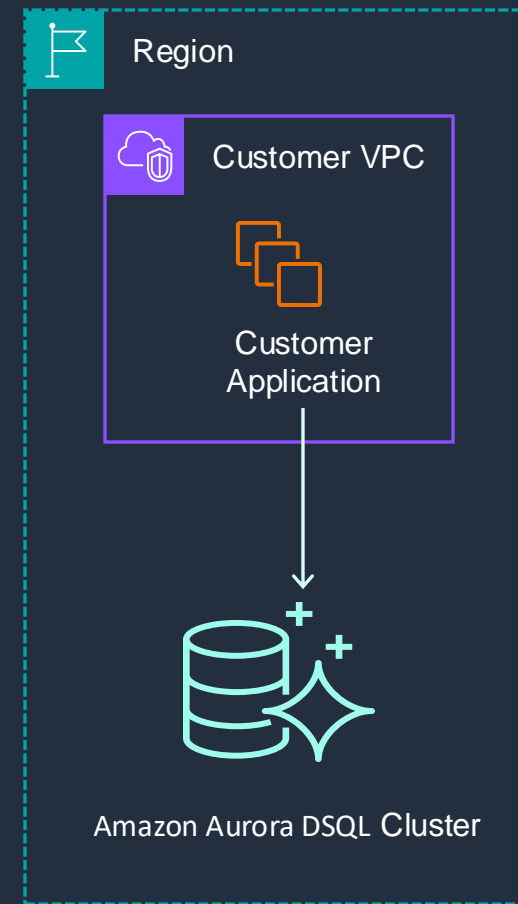
SINGLE REGION CLUSTER

Multi-AZ active active

Designed for 99.99% database availability

Active-active data replication to 3 Availability Zones

Fully automated failure recovery with strong data consistency even during AZ-wide events





Amazon Aurora DSQL

MULTI-REGION ACTIVE-ACTIVE CLUSTER

Extend the benefits of single region cluster across two AWS regions


Designed for 99.999% availability

Active-active data access across both regions

Both regions are **equal**. No primary or secondary. Applications access the **same data from any regional end-point**

Fully automated failure recovery with strong data consistency even during **Region-wide events**

Cross-region replication with **ACID** transactions





Amazon Aurora DSQL

MULTI-REGION ACTIVE-ACTIVE CLUSTER

Extend the benefits of single region cluster across two AWS regions

Active-active data access across both regions

Both regions are **equal**. No primary or secondary. Applications access the **same data** from **any regional end-point**

Fully automated failure recovery with strong data consistency even during Region-wide events

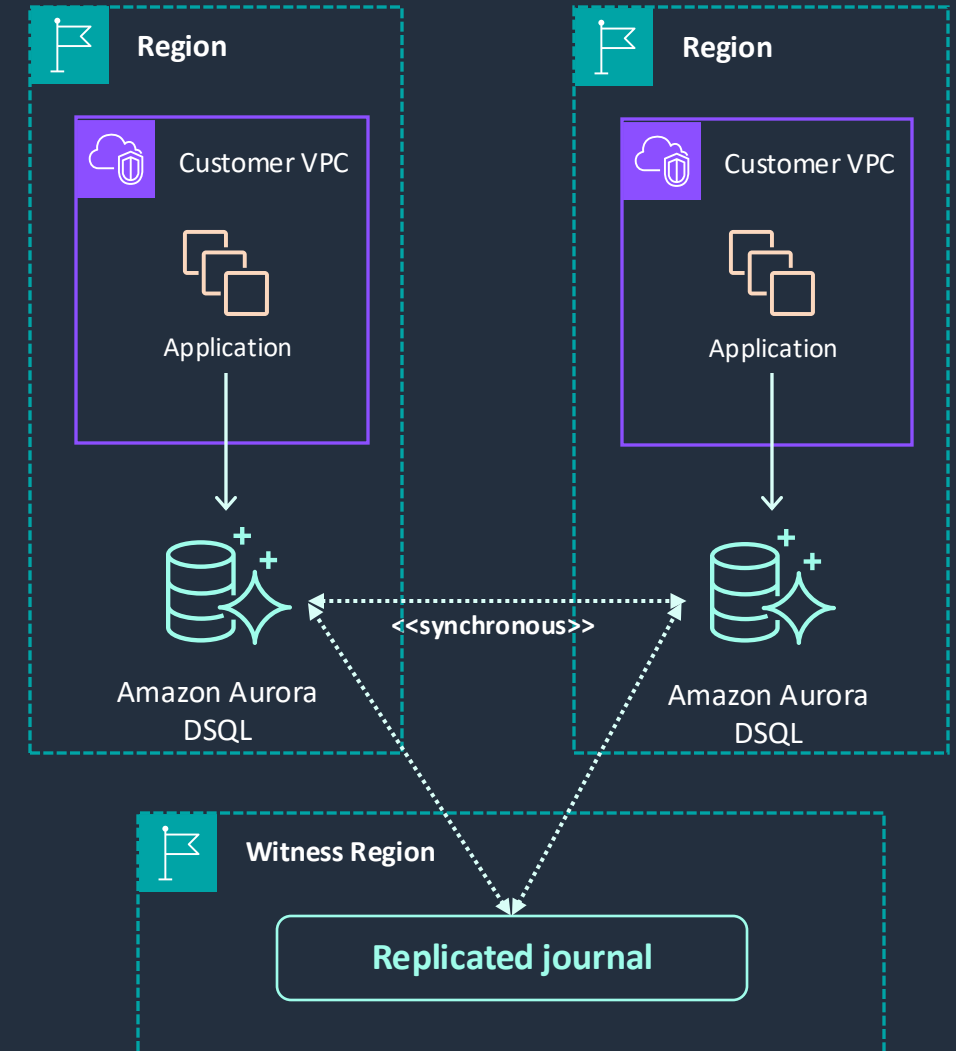
Designed for 99.999% availability

Elastic scaling with no planned downtime

Compute and Storage scale independently

No instances to size, scale, place, or manage

No patching or maintenance windows



Demo





Thank you!

Nethravathi Muddarajaiah

Principal Database Solutions Architect
AWS India

